VEP 115.2 fork scaling on Apple M2 / Docker

HG002 GRCh38 chr22 | 50,861 normalized variants | merged cache 115 | --everything --hgvs

MacBook Air M2 (4P + 4E cores), 24 GiB RAM | native ARM64 Docker, 8 vCPU / 15.35 GiB VM RAM

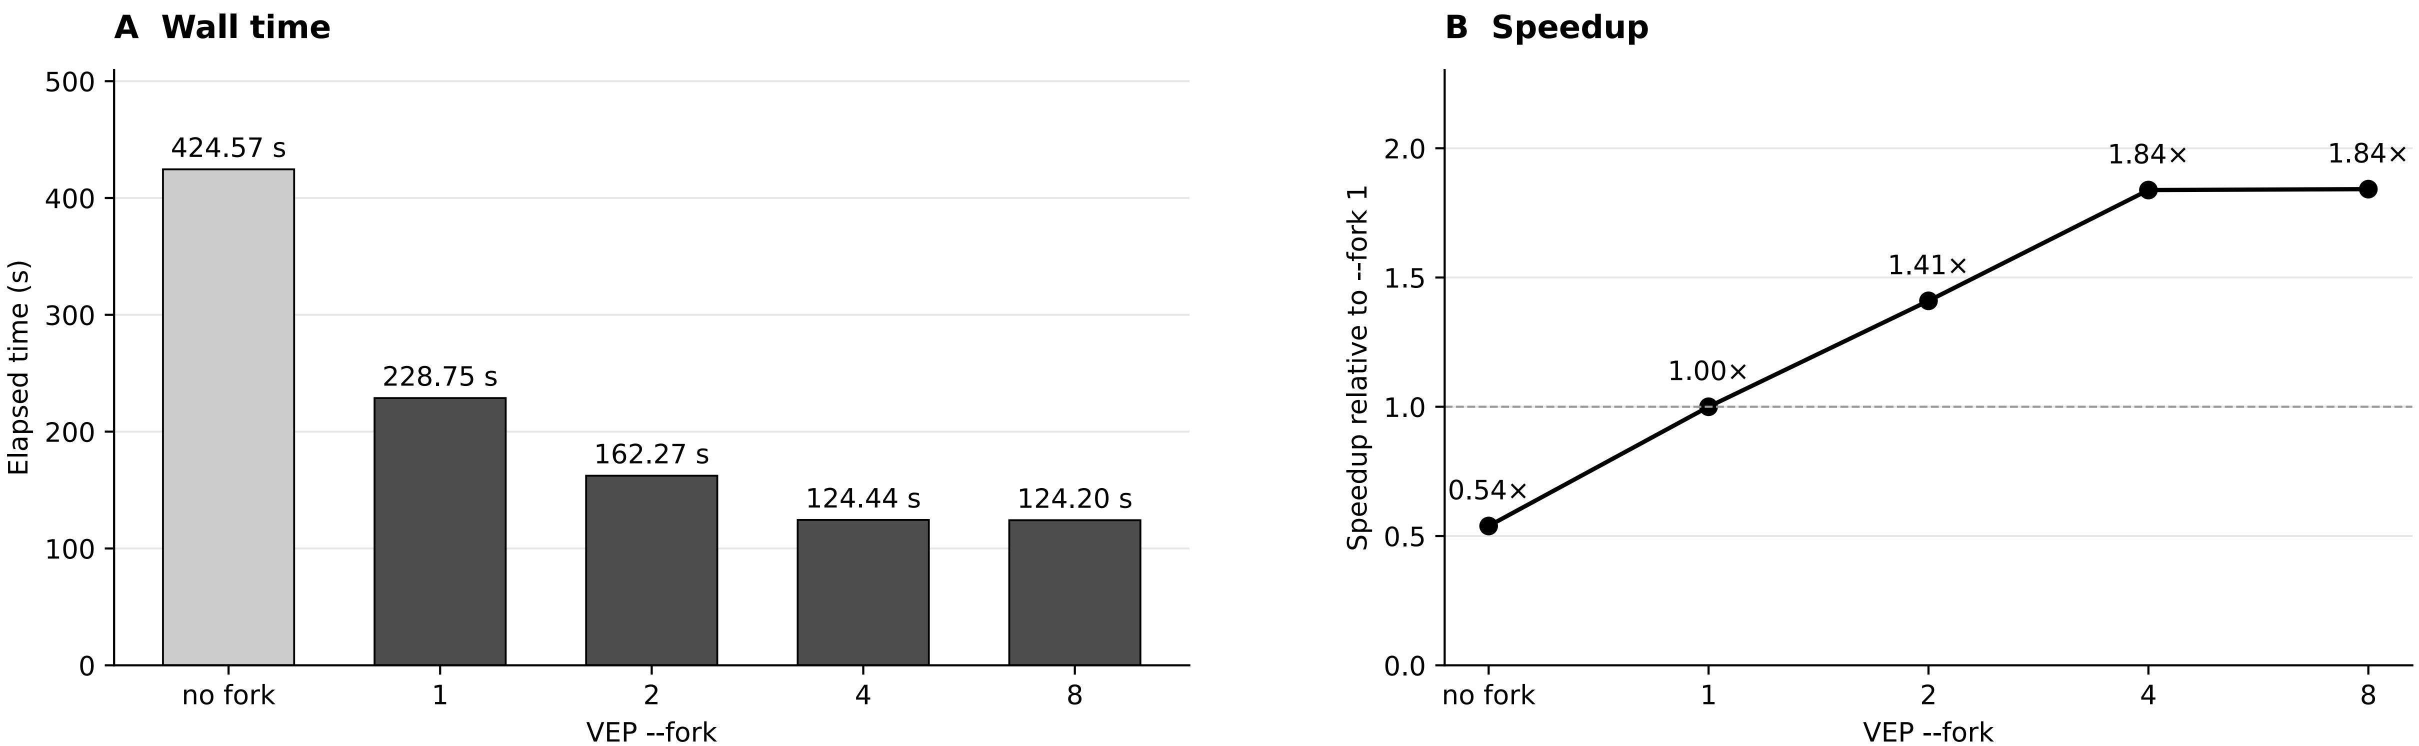

One measured run per setting after a separate --fork 8 warm-up; serial case resumed.

Exact benchmark command

```
cd /Users/tgambin/workspace/vepyr
env \
  DATA_VEPYR_DIR=/Users/tgambin/workspace/data_vepyr \
  INPUT_DIR=/Users/tgambin/workspace/vepyr/performance-tests/vep/outputs/115/macos_m2_chr22_20260907/work/input \
  RELEASE=115 \
  VEP_IMAGE=ensemblorg/ensembl-vep:release_115.2 \
  TIME_BIN=/opt/homebrew/bin/gtime \
  KEEP_VCFS=1 \
  OUT_DIR=/Users/tgambin/workspace/vepyr/performance-tests/vep/outputs/115/macos_m2_chr22_20260907/raw \
  bash performance-tests/vep/scripts/run_vep_fork_scaling.sh merged 1 2 4 8 none

# After a Bash 3.2 compatibility fix, the serial case was resumed with the same
# environment and: bash performance-tests/vep/scripts/run_vep_fork_scaling.sh merged none
```

Clock: GNU time around docker run; startup and plain-VCF writing included.