



VEP 116.0 fork scaling on Apple M3 Max / Docker

HG002 GRCh38 chr22 | 50,861 normalized variants | merged cache 116 | --everything --hgvs

MacBook Pro M3 Max (12P + 4E cores), 64 GiB RAM | native ARM64 Docker Desktop 29.6.2, 16 vCPU / 16.5 GiB VM RAM



One measured run per setting after a separate --fork 8 warm-up.

Exact benchmark command

```
cd /Users/mwiewior/research/git/vepyr
env \
  DATA_VEPYR_DIR=/Users/mwiewior/workspace/data_vepyr \
  INPUT_DIR=/Users/mwiewior/research/git/vepyr/performance-tests/vep/outputs/116/macos_mwiewior_chr22_20260913/work/input \
  RELEASE=116 \
  VEP_IMAGE=ensemblorg/ensembl-vep:release_116.0 \
  TIME_BIN=/opt/homebrew/bin/gtime \
  KEEP_VCFS=1 \
  OUT_DIR=/Users/mwiewior/research/git/vepyr/performance-tests/vep/outputs/116/macos_mwiewior_chr22_20260913/raw \
  bash performance-tests/vep/scripts/run_vep_fork_scaling.sh merged 1 2 4 8 none
```

Clock: GNU time around docker run; startup and plain-VCF writing included.