


VEP 116.0 fork scaling on Apple M3 Max / Docker

HG002 GRCh38 whole genome | 4,096,123 normalized variants | merged cache 116 | --everything --hgvs


MacBook Pro M3 Max (12P + 4E cores), 64 GiB RAM | native ARM64 Docker Desktop 29.6.2, 16 vCPU / 16.5 GiB VM RAM

A Wall time



One measured run per setting, no separate warm-up.

B Speedup



Exact benchmark command

```
cd ~/research/git/vepyr && env INPUT_DIR=/Users/mwiewior/workspace/data_vepyr/input \
  DATA_VEPYR_DIR=/Users/mwiewior/workspace/data_vepyr RELEASE=116 \
  VEP_IMAGE=ensemblorg/ensembl-vep:release_116.0 TIME_BIN=/opt/homebrew/bin/gtime KEEP_VCFS=0 \
  OUT_DIR=/Users/mwiewior/workspace/data_vepyr/output/116/vep_merged_wgs_macos_20260914/raw \
  bash performance-tests/vep/scripts/run_vep_fork_scaling.sh merged 8 4 2 1 none
```

Clock: GNU time around docker run; startup and plain-VCF writing included.